

FPC lista 1 - questão 5: boxplots

David Ferreira Quaresma (david.quaresma@ccc.ufcg.edu.br), Renato Dantas Henriques
(renato.henriques@ccc.ufcg.edu.br)

junho, 2019

Letra A

```
concurrentHash = read.csv("./results/concurrentHash-t1024-r0.5.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1024-r0.5.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 8500 13000 | 50: 476 482 | 95: 1656 1715 |
## 99: 3674 4202 | 99.9: 4062000 6602000 | 99.99: 10600000 17580000 | Dist.Tail.: 10600000 17580000
## Latencia(ns) synchronizedMap avg: 580 3500 | 50: 609 616 | 95: 1922 1992 |
## 99: 3997 4293 | 99.9: 58610 173100 | 99.99: 917000 3184000 | Dist.Tail.: 916400 3183000
```

ECDF - 1024 threads, 50/50 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1024-r0.5.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1024-r0.5.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 8500 13000 | 50: 476 482 | 95: 1656 1715 |
## 99: 3674 4202 | 99.9: 4062000 6602000 | 99.99: 10600000 17580000 | Dist.Tail.: 10600000 17580000
## Latencia(ns) synchronizedMap avg: 580 3500 | 50: 609 616 | 95: 1922 1992 |
## 99: 3997 4293 | 99.9: 58610 173100 | 99.99: 917000 3184000 | Dist.Tail.: 916400 3183000
```

ECDF - 1024 threads, 90/10 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1024-r0.9.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1024-r0.9.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 16000 42000 | 50: 500 505 | 95: 1499 1537 |
## 99: 3602 3989 | 99.9: 27460 207800 | 99.99: 74880000 172400000 | Dist.Tail.: 74880000 172400000
## Latencia(ns) synchronizedMap avg: 1000 1300 | 50: 603 611 | 95: 2003 2055 |
## 99: 3551 3927 | 99.9: 26070 148300 | 99.99: 469700 1206000 | Dist.Tail.: 469100 1205000
```

ECDF - 1024 threads, 10/90 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1024-r0.1.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1024-r0.1.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 6900 13000 | 50: 291 296 | 95: 1612 1667 |
## 99: 3321 3627 | 99.9: 127400 965900 | 99.99: 21380000 21460000 | Dist.Tail.: 21380000 21460000
## Latencia(ns) synchronizedMap avg: 1200 1900 | 50: 379 383 | 95: 2326 2393 |
## 99: 4433 4841 | 99.9: 45710 175300 | 99.99: 1196000 3388000 | Dist.Tail.: 1196000 3388000
```

ECDF - 128 threads, 50/50 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t128-r0.5.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t128-r0.5.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 7000 13000 | 50: 553 559 | 95: 1696 1739 |
## 99: 3281 3618 | 99.9: 166400 1559000 | 99.99: 15450000 26230000 | Dist.Tail.: 15450000 26230000
## Latencia(ns) synchronizedMap avg: 910 1300 | 50: 593 600 | 95: 1954 2009 |
## 99: 3623 3817 | 99.9: 14230 27840 | 99.99: 242400 1857000 | Dist.Tail.: 241800 1856000
```

ECDF - 128 threads, 90/10 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t128-r0.9.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t128-r0.9.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 2300 4600 | 50: 560 564 | 95: 1553 1589 |
## 99: 2781 2989 | 99.9: 21600 178300 | 99.99: 9745000 9990000 | Dist.Tail.: 9744000 9990000
## Latencia(ns) synchronizedMap avg: 930 1100 | 50: 612 618 | 95: 1835 1873 |
## 99: 3237 3448 | 99.9: 15710 23980 | 99.99: 260400 924200 | Dist.Tail.: 259800 923500
```

ECDF - 128 threads, 10/90 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t128-r0.1.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t128-r0.1.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 4400 8700 | 50: 421 426 | 95: 1716 1766 |
## 99: 3583 3932 | 99.9: 144900 223200 | 99.99: 8113000 22790000 | Dist.Tail.: 8113000 22790000
## Latencia(ns) synchronizedMap avg: 820 1400 | 50: 366 371 | 95: 2093 2150 |
## 99: 3922 4289 | 99.9: 17040 26570 | 99.99: 227900 2781000 | Dist.Tail.: 227500 2780000
```

ECDF - 1 threads, 50/50 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1-r0.5.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1-r0.5.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 540 1100 | 50: 387 389 | 95: 667 681 |
## 99: 1703 1867 | 99.9: 11700 18910 | 99.99: 198500 4010000 | Dist.Tail.: 198200 4010000
## Latencia(ns) synchronizedMap avg: 580 700 | 50: 352 355 | 95: 999 1016 |
## 99: 2057 2147 | 99.9: 11780 21330 | 99.99: 193500 603500 | Dist.Tail.: 193100 603200
```

ECDF - 1 threads, 90/10 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1-r0.9.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1-r0.9.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 420 670 | 50: 290 293 | 95: 629 641 |
## 99: 1235 1406 | 99.9: 11480 13850 | 99.99: 190100 1462000 | Dist.Tail.: 189800 1462000
## Latencia(ns) synchronizedMap avg: 420 480 | 50: 304 308 | 95: 637 648 |
## 99: 1415 1509 | 99.9: 11050 12200 | 99.99: 167900 195700 | Dist.Tail.: 167600 195400
```

ECDF - 1 threads, 10/90 read/write

```
concurrentHash = read.csv("./results/concurrentHash-t1-r0.1.csv", header=T, dec=".")
synchronizedMap = read.csv("./results/synchronizedMap-t1-r0.1.csv", header=T, dec=".")
concurrentHash <- filter(concurrentHash, time_in_nanoseconds > 0)
synchronizedMap <- filter(synchronizedMap, time_in_nanoseconds > 0)
print_summary_table(concurrentHash, synchronizedMap, "concurrentHash", "synchronizedMap")
```

```
## Latencia(ns) concurrentHash avg: 510 1100 | 50: 250 252 | 95: 665 691 |
## 99: 1823 1883 | 99.9: 11140 12280 | 99.99: 262700 4458000 | Dist.Tail.: 262500 4458000
## Latencia(ns) synchronizedMap avg: 490 880 | 50: 271 274 | 95: 966.9 989 |
## 99: 2260 2358 | 99.9: 10790 18700 | 99.99: 191800 2074000 | Dist.Tail.: 191500 2073000
```

Letra B

ECDF - 1024 threads, 50/50 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1024-r0.5.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1024-r0.5.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")

## Latencia(ns) copyOnWriteList avg: 3200000 3500000 | 50: 1003 3790 | 95: 18890000 22480000 |
## 99: 86810000 92110000 | 99.9: 129200000 134100000 | 99.99: 146400000 151100000 | Dist.Tail.: 146400000
## Latencia(ns) synchronizedList avg: 540 660 | 50: 282 285 | 95: 1252 1296 |
## 99: 2358 2503 | 99.9: 13630 37280 | 99.99: 203600 565900 | Dist.Tail.: 203400 565600
```

ECDF - 1024 threads, 90/10 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1024-r0.9.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1024-r0.9.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")

## Latencia(ns) copyOnWriteList avg: 1800 2300 | 50: 138 140 | 95: 9685 10280 |
## 99: 23120 24130 | 99.9: 102200 193400 | 99.99: 931100 1830000 | Dist.Tail.: 931000 1830000
## Latencia(ns) synchronizedList avg: 360 690 | 50: 214 216 | 95: 856.3 890 |
## 99: 2093 2219 | 99.9: 8276 17210 | 99.99: 190400 1255000 | Dist.Tail.: 190200 1255000
```

ECDF - 1024 threads, 10/90 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1024-r0.1.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1024-r0.1.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")

## Latencia(ns) copyOnWriteList avg: 3.7e+07 3.8e+07 | 50: 58090 59470 | 95: 1.93e+08 197200000 |
## 99: 326500000 347800000 | 99.9: 475400000 508800000 | 99.99: 546100000 554700000 | Dist.Tail.: 5.4e+08
## Latencia(ns) synchronizedList avg: 310 1600 | 50: 140 141 | 95: 982 1030 |
## 99: 2048 2187 | 99.9: 8430 13500 | 99.99: 215700 12460000 | Dist.Tail.: 215600 12460000
```

ECDF - 128 threads, 50/50 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t128-r0.5.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t128-r0.5.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 1e+05 140000 | 50: 776 2004 | 95: 65760 69330 |
## 99: 400800 509100 | 99.9: 24290000 50930000 | 99.99: 77400000 90570000 | Dist.Tail.: 77400000 905
## Latencia(ns) synchronizedList avg: 510 700 | 50: 315 318 | 95: 1118 1147 |
## 99: 1916 2007 | 99.9: 7047 12960 | 99.99: 176300 1390000 | Dist.Tail.: 176000 1390000
```

ECDF - 128 threads, 90/10 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t128-r0.9.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t128-r0.9.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 2100 3600 | 50: 216 218 | 95: 11840 12740 |
## 99: 24980 26010 | 99.9: 44820 177400 | 99.99: 950200 11380000 | Dist.Tail.: 949900 11380000
## Latencia(ns) synchronizedList avg: 420 490 | 50: 284 286 | 95: 961 982 |
## 99: 1741 1815 | 99.9: 5605 10480 | 99.99: 176600 236800 | Dist.Tail.: 176300 236500
```

ECDF - 128 threads, 10/90 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t128-r0.1.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t128-r0.1.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 5e+06 5200000 | 50: 62490 63620 | 95: 31420000 32550000 |
## 99: 59940000 62270000 | 99.9: 103200000 116600000 | 99.99: 129300000 144900000 | Dist.Tail.: 1292
## Latencia(ns) synchronizedList avg: 580 850 | 50: 343 345 | 95: 1332 1375 |
## 99: 2113 2222 | 99.9: 11200 19490 | 99.99: 248800 1080000 | Dist.Tail.: 248500 1080000
```

ECDF - 1 threads, 50/50 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1-r0.5.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1-r0.5.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 13000 15000 | 50: 686 1657 | 95: 35010 36170 |
## 99: 72550 74060 | 99.9: 457300 705600 | 99.99: 4069000 10040000 | Dist.Tail.: 4068000 10040000
## Latencia(ns) synchronizedList avg: 240 280 | 50: 177 179 | 95: 352 357 |
## 99: 958 1082 | 99.9: 6747 10760 | 99.99: 124700 189500 | Dist.Tail.: 124500 189300
```

ECDF - 1 threads, 90/10 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1-r0.9.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1-r0.9.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 980 1100 | 50: 145 146 | 95: 7875 8813 |
## 99: 13370 13770 | 99.9: 24880 31640 | 99.99: 205200 406900 | Dist.Tail.: 205100 406800
## Latencia(ns) synchronizedList avg: 210 240 | 50: 141 141 | 95: 317 326 |
## 99: 1071 1213 | 99.9: 2439 5609 | 99.99: 149100 166700 | Dist.Tail.: 149000 166500
```

ECDF - 1 threads, 10/90 read/write

```
copyOnWriteList = read.csv("./results/copyOnWriteList-t1-r0.1.csv", header=T, dec=".")
synchronizedList = read.csv("./results/synchronizedList-t1-r0.1.csv", header=T, dec=".")
copyOnWriteList <- filter(copyOnWriteList, time_in_nanoseconds > 0)
synchronizedList <- filter(synchronizedList, time_in_nanoseconds > 0)
print_summary_table(copyOnWriteList, synchronizedList, "copyOnWriteList", "synchronizedList")
```

```
## Latencia(ns) copyOnWriteList avg: 39000 42000 | 50: 27800 28370 | 95: 82390 86290 |
## 99: 172100 213000 | 99.9: 896600 1146000 | 99.99: 6482000 9859000 | Dist.Tail.: 6455000 9830000
## Latencia(ns) synchronizedList avg: 230 270 | 50: 155 155 | 95: 391 397 |
## 99: 834.8 881 | 99.9: 2329 7104 | 99.99: 43600 314500 | Dist.Tail.: 43450 314300
```